

```
In [1]: # Importing required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import joblib

# Load the dataset
link = 'https://github.com/dsrscientist/DSData/raw/master/loan_prediction.csv'
data = pd.read_csv(link)

# EDA Analysis
# Display basic statistics
print(data.describe())

# Check for missing values
print(data.isnull().sum())

# Visualize the distribution of the target variable (Loan_Status)
plt.figure(figsize=(6, 4))
sns.countplot(x='Loan_Status', data=data)
plt.title('Distribution of Loan Status')
plt.show()

# Preprocessing and Feature Engineering
# Handling missing values
data['Gender'].fillna(data['Gender'].mode()[0], inplace=True)
data['Married'].fillna(data['Married'].mode()[0], inplace=True)
data['Dependents'].fillna(data['Dependents'].mode()[0], inplace=True)
data['Self_Employed'].fillna(data['Self_Employed'].mode()[0], inplace=True)
data['LoanAmount'].fillna(data['LoanAmount'].median(), inplace=True)
data['Loan_Amount_Term'].fillna(data['Loan_Amount_Term'].mode()[0], inplace=True)
data['Credit_History'].fillna(data['Credit_History'].mode()[0], inplace=True)

# Encoding categorical variables
le = LabelEncoder()
data['Gender'] = le.fit_transform(data['Gender'])
data['Married'] = le.fit_transform(data['Married'])
data['Dependents'] = le.fit_transform(data['Dependents'])
data['Education'] = le.fit_transform(data['Education'])
data['Self_Employed'] = le.fit_transform(data['Self_Employed'])
data['Property_Area'] = le.fit_transform(data['Property_Area'])
data['Loan_Status'] = le.fit_transform(data['Loan_Status'])

# Feature selection
features = ['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed',
            'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History', 'Property_Area']
target = 'Loan_Status'

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(data[features], data[target], test_size=0.2, random_state=42)

# Build/Test Multiple Models
# Random Forest Classifier
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)

# Check for overfitting/underfitting
train_predictions = rf_model.predict(X_train)
print(f"Train Accuracy: {accuracy_score(y_train, train_predictions)}")
print(f"Test Accuracy: {accuracy_score(y_test, rf_predictions)}")

# Cross-validation
cv_accuracy = cross_val_score(rf_model, data[features], data[target], scoring='accuracy', cv=5)
print(f"Cross-Validation Accuracy: {np.mean(cv_accuracy)}")

# Hyperparameter Tuning
param_grid = {'n_estimators': [50, 100, 200], 'max_depth': [None, 10, 20], 'min_samples_split': [2, 5, 10]}
grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid, scoring='accuracy', cv=5)
grid_search.fit(data[features], data[target])

# Select the Best Model
best_model = grid_search.best_estimator_

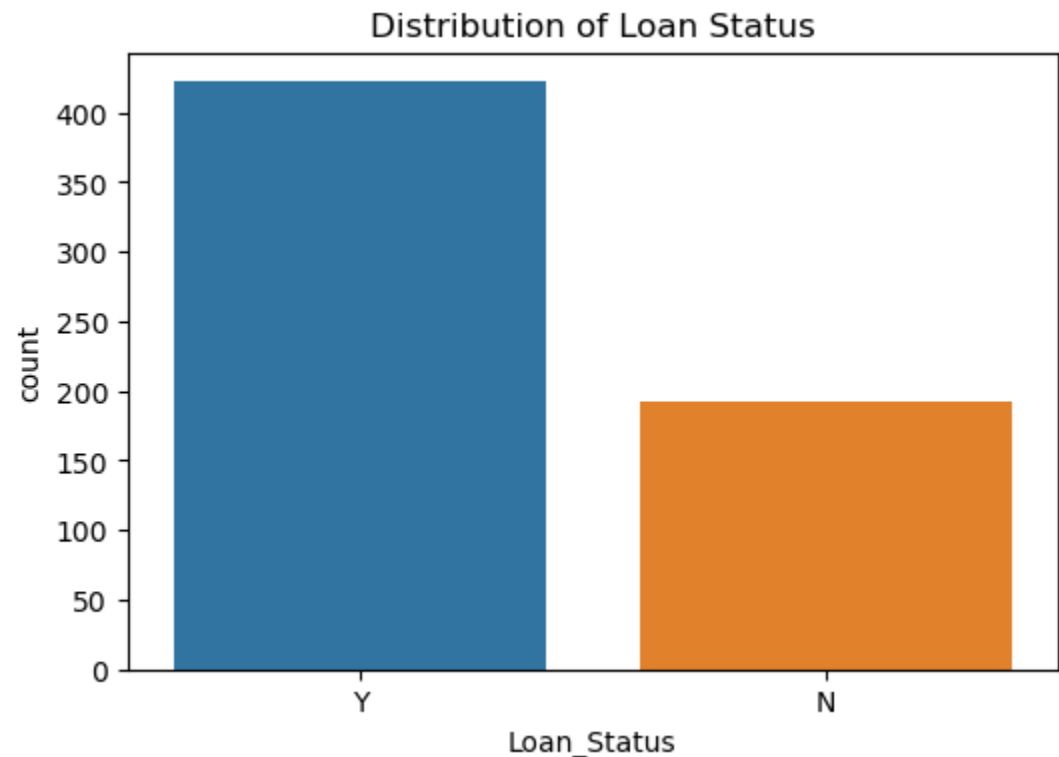
# Save the Best Model for Production
joblib.dump(best_model, 'loan_approval_model.pkl')

# Explanation for Model Selection
# The best model is selected based on the highest accuracy obtained during cross-validation.
# This metric provides a measure of how well the model performs in predicting loan approval status.

# Additional Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, rf_predictions))
print("Confusion Matrix:\n", confusion_matrix(y_test, rf_predictions))
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
count	614.000000	614.000000	592.000000	600.00000	
mean	5403.459283	1621.245798	146.412162	342.00000	
std	6109.041673	2926.248369	85.587325	65.12041	
min	150.000000	0.000000	9.000000	12.00000	
25%	2877.500000	0.000000	100.000000	360.00000	
50%	3812.500000	1188.500000	128.000000	360.00000	
75%	5795.000000	2297.250000	168.000000	360.00000	
max	81000.000000	41667.000000	700.000000	480.00000	

	Credit_History
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000
Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0
dtype:	int64



Train Accuracy: 1.0				
Test Accuracy: 0.7723577235772358				
Cross-Validation Accuracy: 0.7882713581234173				
Classification Report:				
	precision	recall	f1-score	support
0	0.86	0.42	0.56	43
1	0.75	0.96	0.85	80
accuracy			0.77	123
macro avg	0.81	0.69	0.70	123
weighted avg	0.79	0.77	0.75	123
Confusion Matrix:				
[[18 25]				
[ 3 77]]				

Importing Libraries: You start by importing the necessary Python libraries for data analysis, visualization, and machine learning.

Loading the Dataset: You load a dataset related to loan predictions from a given link using pandas.

Exploratory Data Analysis (EDA): You conduct exploratory data analysis by displaying basic statistics, checking for missing values, and visualizing the distribution of the target variable (Loan\_Status).

Preprocessing and Feature Engineering: You handle missing values by imputing them with mode or median values. Categorical variables are encoded using LabelEncoder. Additionally, you select specific features and the target variable for your model.

Train-Test Split: You split the dataset into training and testing sets using the train\_test\_split function.

Building and Testing Multiple Models: You use a Random Forest Classifier, fit the model, and make predictions. You check for overfitting/underfitting by comparing the accuracy on both training and testing sets. Cross-validation is performed to obtain a more robust evaluation.

Hyperparameter Tuning: You perform a grid search to find the best hyperparameters for the Random Forest model.

Selecting the Best Model: The best model is selected based on the results of the hyperparameter tuning.

Saving the Model: The best model is saved using the joblib library for potential use in production.

Explanation for Model Selection: A brief explanation is provided on how the best model is selected based on cross-validation accuracy.

Additional Evaluation Metrics: You print the classification report and confusion matrix for further evaluation of the model's performance.